Artificial Intelligence
CS-401

Chapter # 03
Solving Problems by Searching – Informed Searches

Dr. Hafeez Ur Rehman
(Email: hafeez.urrehman@nu.edu.pk)

Today's Outline

- Limitations of uninformed search methods

- **Informed (or heuristic)** search uses problem-specific **heuristics** to improve efficiency
  - Greedy Best-First Search
  - A* Search
  - RBFS (Recursive Best First Search)
  - SMA* (Simplified Memory Bounded A*)

- Can provide significant speed-ups in practice
  - e.g., on 8-puzzle
  - But can still have worst-case exponential time complexity

Limitations of uninformed search

- For 8-puzzle
  - Avg. solution cost is about 22 steps
  - branching factor ~ 3
  - Exhaustive search to depth 22:
    - **$3^{22} = 3.1 \times 10^{10}$ states**
  - For example, at d=12, IDS expands 3.6 million states on average
  - If we slightly complicate the problem to a 5x5 (24 puzzle) then the state space grows to **10^{24}** states. For such a problem, uninformed search strategies are the wrost choice.

Start State   Goal State

Best-first search Class of algorithms

- Idea: use an evaluation function f(n) for each node
  - → estimate of "**desirability**"
  - → Expand most desirable unexpanded node

- Implementation:
  - Order the nodes in fringe (queue) by f(n) (i.e. by desirability, highest f(n) first)

- **Special cases:**
  1. **Uniform cost search** (from previous lecture): f(n) = g(n) = path to n
  2. **Greedy best-first search**
  3. **A* search**

- Note: evaluation function f(n) is an <u>estimate</u> of node quality

Heuristic function (for informed searches)

- Heuristic:
  - Definition: "using rules of thumb to find answers". These are clues in the form of additional information to help algorithm find solution quickly.

- Heuristic function h(n)
  1. **Estimate of (optimal) cost from n to goal**
  2. h(n) = 0 if n is a goal node
  3. **Example: straight line distance from n to Bucharest**
     - Note that this is not the true state-space distance
     - It is an estimate – actual state-space distance can be higher
  4. Provides problem-specific knowledge to the search algorithm

1. Greedy best-first search

- Special case of best-first search
  - Uses h(n) = heuristic function as its evaluation function
  - Expand the node that appears closest to goal
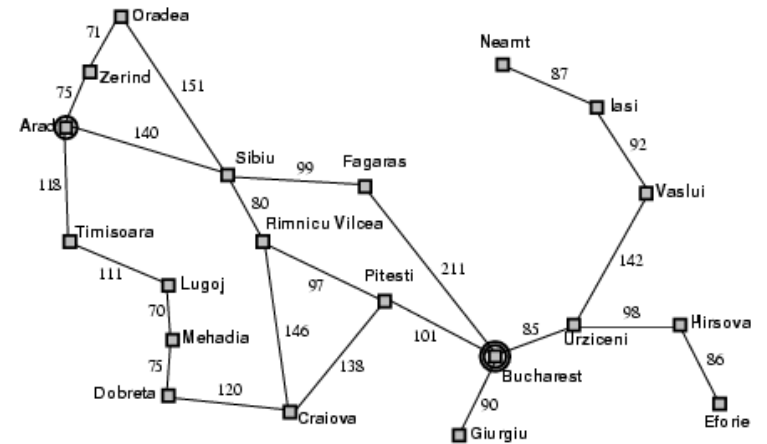
Romania with step costs in km

h(n)

Straight–line distance to Bucharest

|	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search example

# Greedy best-first search example

# Greedy best-first search example

Greedy best-first search example
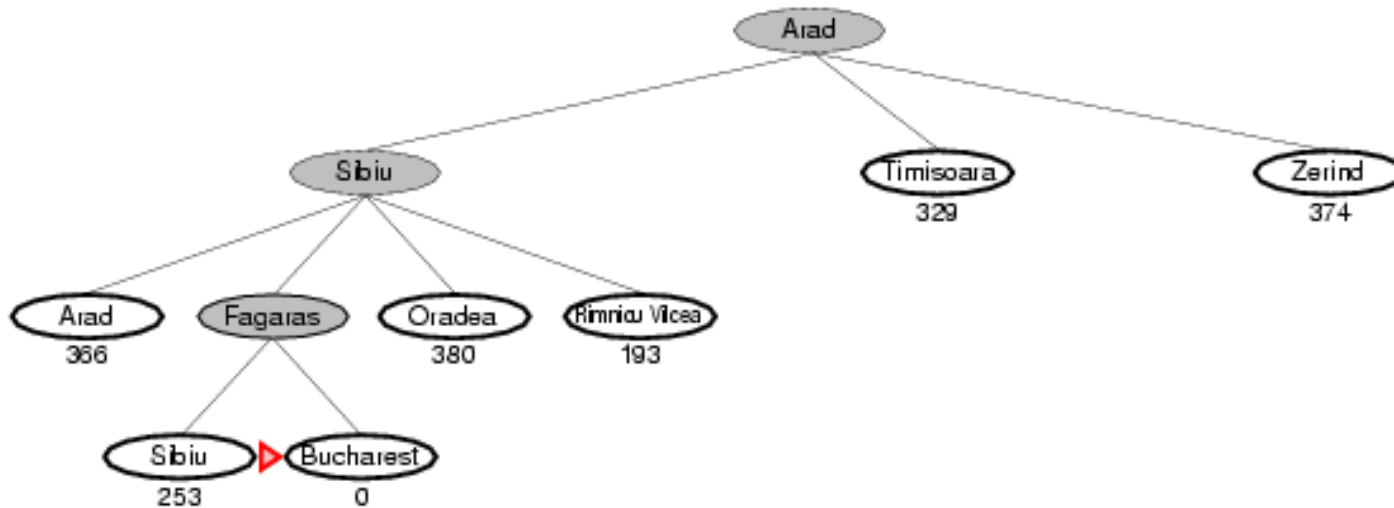
Attention: This is the non-optimal path

BUT Optimal Path

Conclusion: The Greedy Best First Search is non Optimal.

Completeness of Greedy BFS

Figure 3.8: A graph that is bad for best–first search

Example 3.14: Consider the graph shown in Figure 3.8, where the cost of an arc is its length. The aim is to find the shortest path from *s* to *g*. Suppose the Euclidean distance to the goal *g* is used as the heuristic function. A heuristic depth–first search will select the node below *s* and will never terminate. Similarly, because all of the nodes below *s* look good, a best–first search will cycle between them, never trying an alternate route from *s*.

Properties of greedy best-first search

- **Complete?**
  - **No.** unless it keeps track of all states visited
    - Otherwise can get stuck in loops (just like DFS)

- **Optimal?**
  - **No.** We just saw a counter-example (Romania map)

- **Time?**
  - **Exponential** i.e., $O(b^m)$, can generate all nodes at depth m before finding solution
 - Can be much better with a good heuristic

- **Space?**
  - **Exponential** i.e., $O(b^m)$ – again, worst case, can generate all nodes at depth m before finding solution

2. A* Search

- Expand node based on estimate of **total path cost** through node

- Evaluation function $f(n) = g(n) + h(n)$
  - $g(n)$ = cost so far to reach n
  - $h(n)$ = estimated cost from n to goal
  - $f(n)$ = estimated total cost of path through n to goal

- Efficiency of search will depend on **quality of heuristic h(n)**
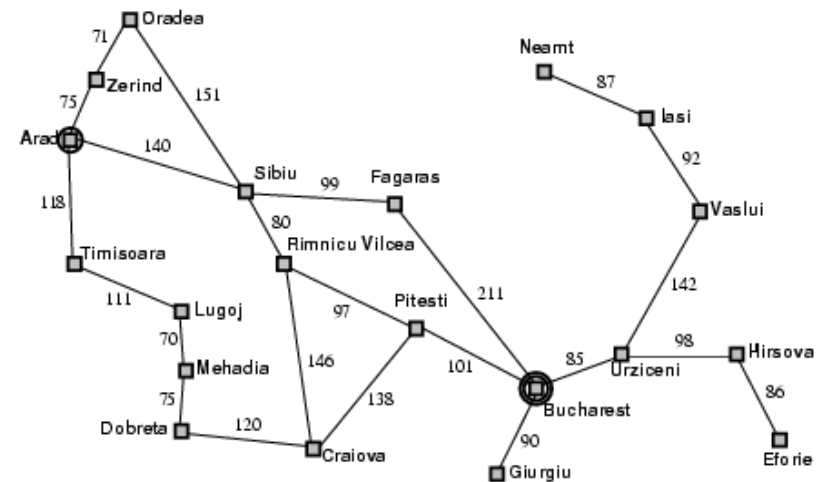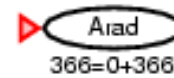
# A* search example [Arad to Bucharest]
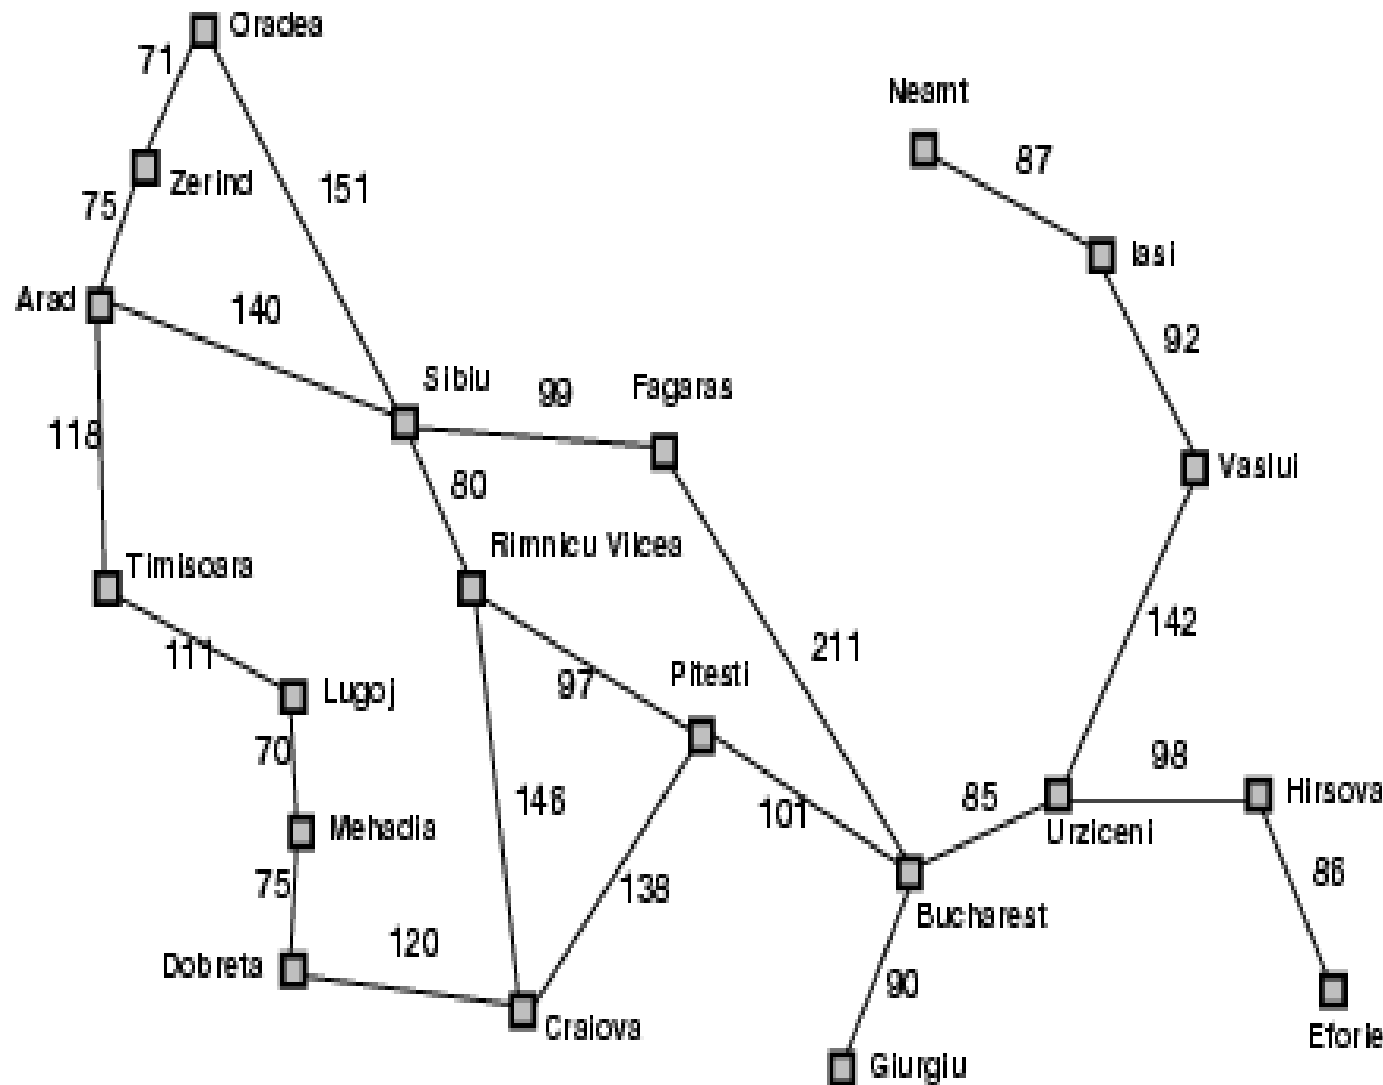
h(n)

Straight—line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

A* search example [Arad to Bucharest]

Arad
366=0+366

A* search example [Arad to Bucharest]

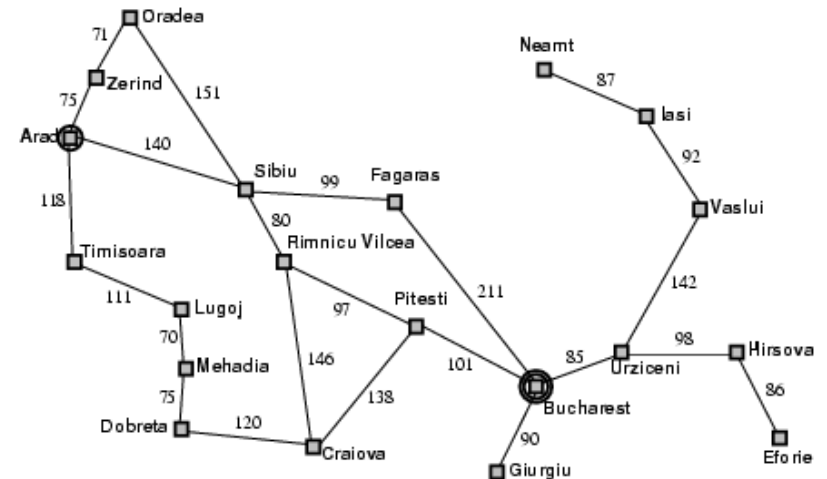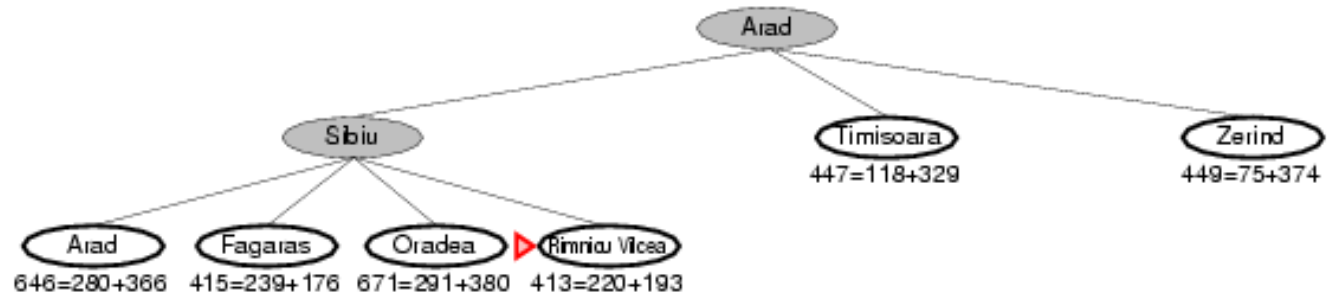# A* search example [Arad to Bucharest]

A* search example [Arad to Bucharest]

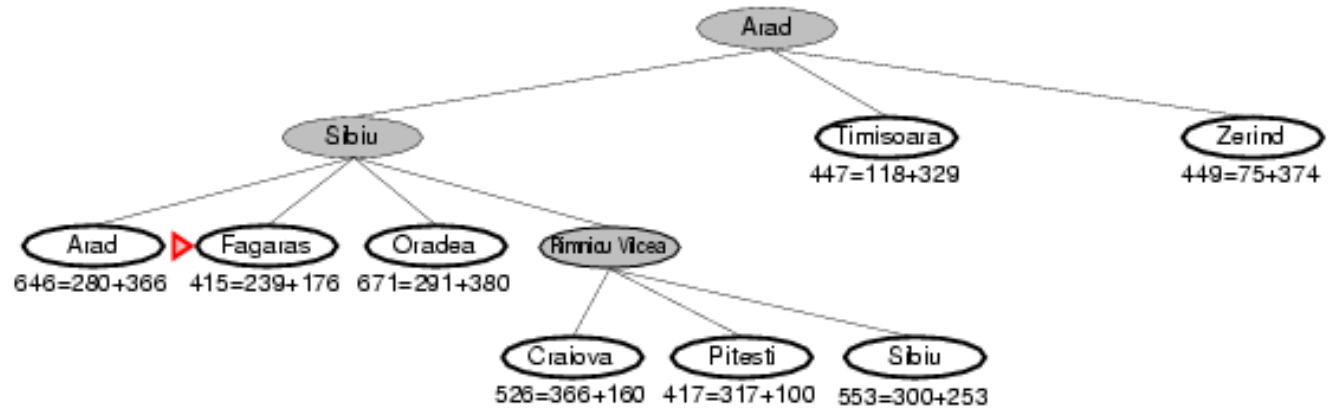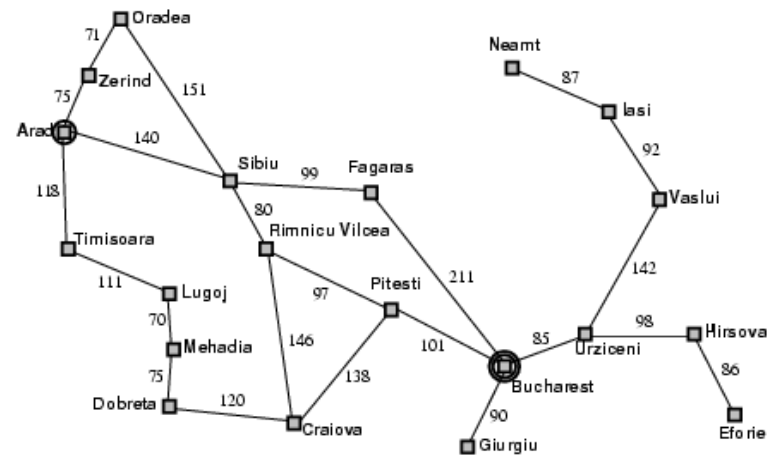# A* search example [Arad to Bucharest]
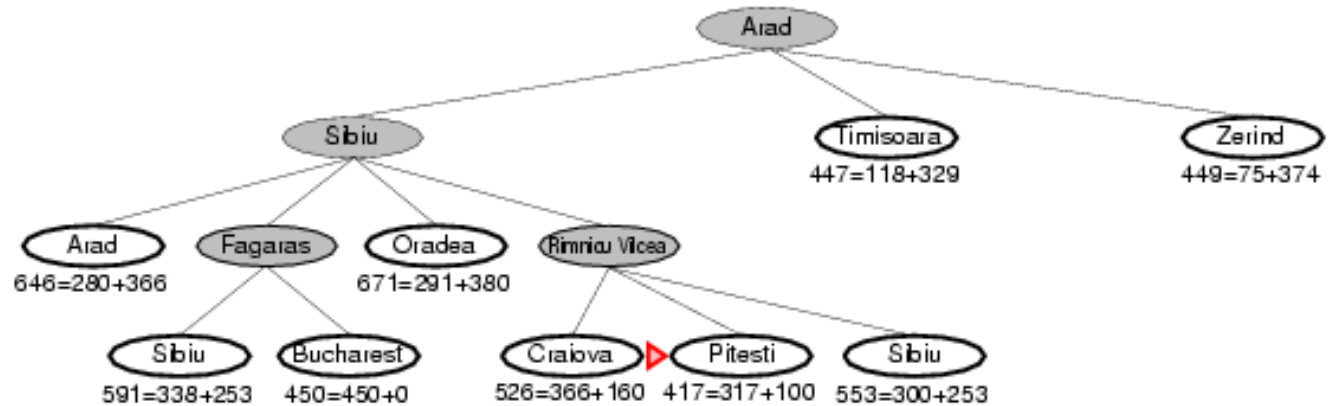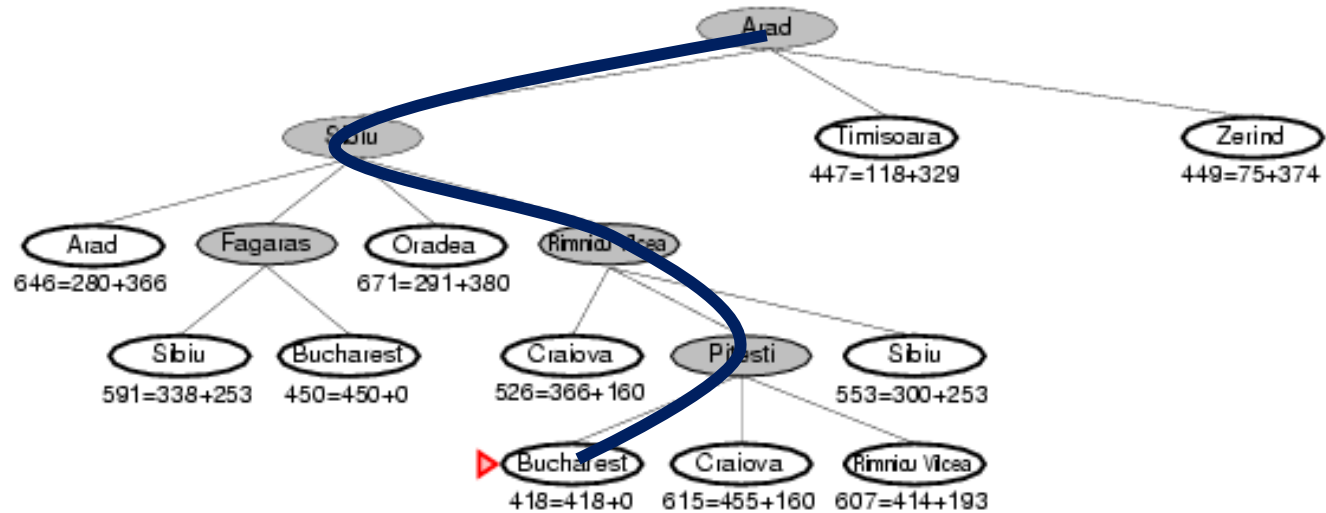
A* search example [Arad to Bucharest]

Finds Optimal Path

But Optimality depends on what???

Admissible heuristics

- A heuristic $h(n)$ is **admissible** if for every node n, $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n.

- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**

- Example: $h_{SLD}(n)$ *is admissible*
  - never overestimates the actual road distance

- Theorem:
  If $h(n)$ is admissible, A* using TREE-SEARCH is optimal

Optimality of A* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G.

- **Need to prove that**
- $f(G_2) > f(G) > f(n)$
- and A* will never select G_2 for expansion

- $f(G_2)$ = $g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G)$ = $g(G)$ since $h(G) = 0$
- $\mathbf{f(G_2)\ >\ f(G)}$ from above

Optimality of A* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G.



- $f(G_2)$ $> f(G)$ from above
- $h(n)$ $\leq h^*(n)$ **since h is admissible**
- $g(n) + h(n)$ $\leq g(n) + h^*(n)$
- $f(n)$ $\leq f(G)$

Hence $f(G_2) > f(G) > f(n)$, and A* will never select G_2 for expansion

Optimality for graphs?

- Admissibility is not sufficient for graph search
  - In graph search, the optimal path to a repeated state could be discarded if it is not the first one generated
  - Can fix problem by requiring <u>consistency property</u> for h(n)

- A heuristic is consistent if for every successor *n'* of a node *n* generated by any action *a*,

  $$h(n) \leq c(n,a,n') + h(n')$$

  *(**aka "monotonic"**)*
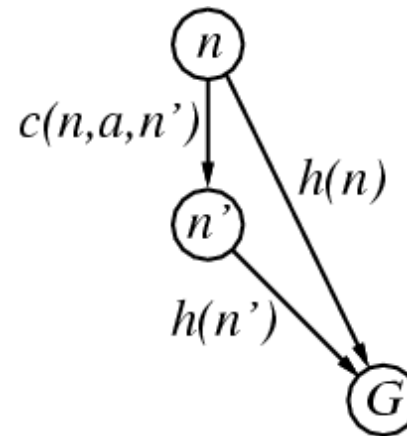
- admissible heuristics are generally consistent

A* is optimal with consistent heuristics

- If *h* is consistent, we have

$$f(n') = g(n') + h(n')$$
$$= g(n) + c(n,a,n') + h(n')$$
$$\geq g(n) + h(n)$$
$$\geq f(n)$$

  i.e., *f(n)* is non-decreasing along any path.

  Thus, first goal-state selected for expansion must be optimal

- Theorem:
  - If *h(n)* is consistent, A* using `GRAPH-SEARCH` is optimal

Contours of A* Search

- A* expands nodes in order of increasing f value
- Gradually adds "f-contours" of nodes
- Contour i has all nodes with $f=f_i$, where $f_i < f_{i+1}$
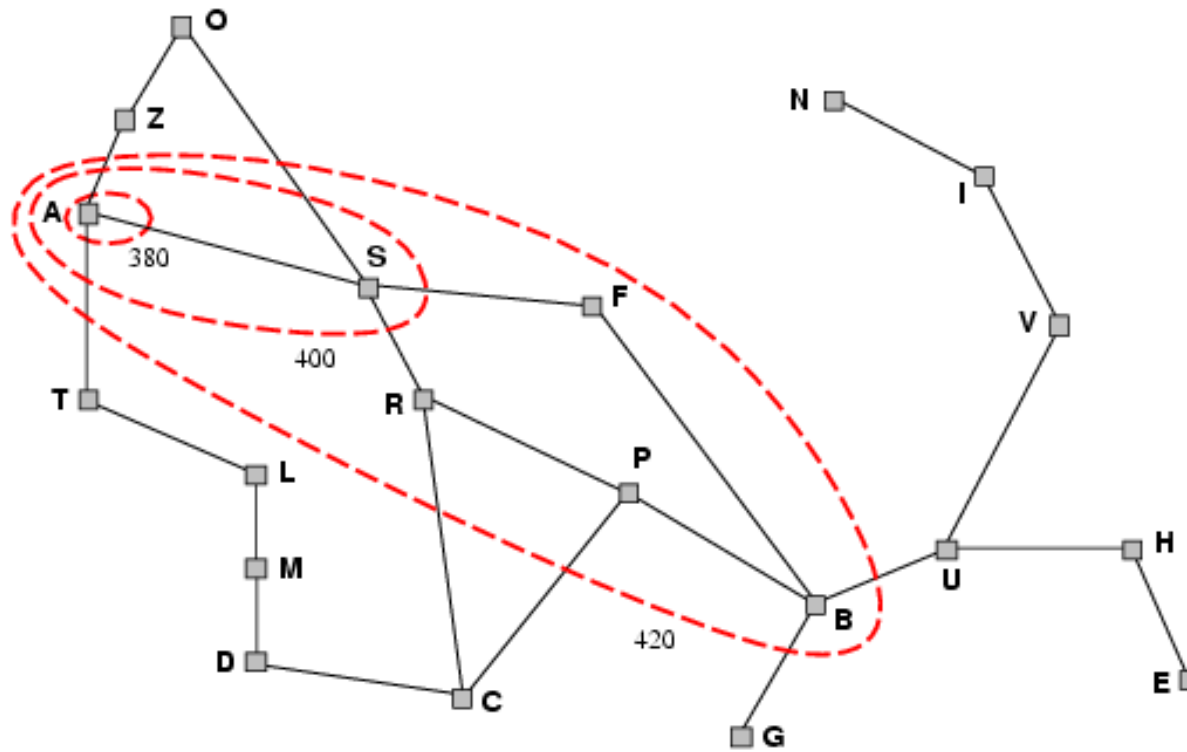
Contours of A* Search
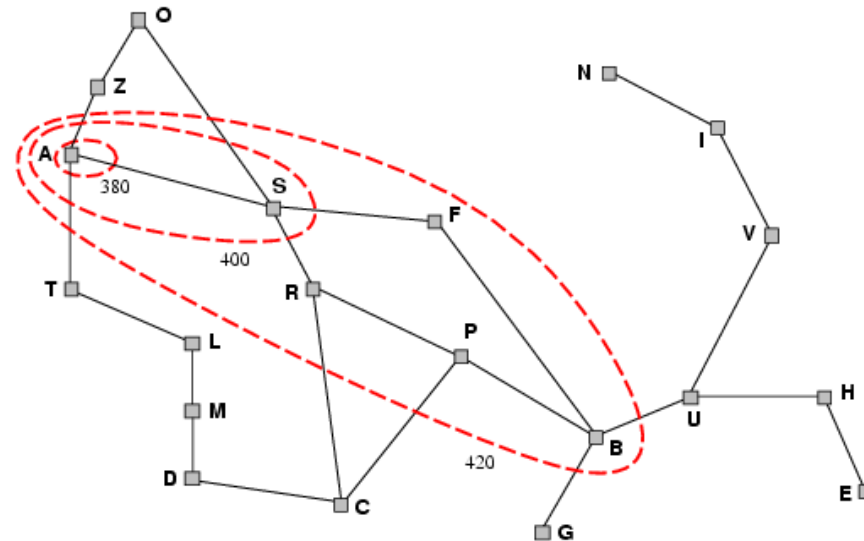
- With uniform-cost h(n) = 0, contours will be circular

- A* with good heuristics, contours will be focused around optimal path

- A* will expand all nodes with cost f(n) < C*

Properties of A*

- Complete?
  - **Yes** (unless there are infinitely many nodes with f ≤ *f(G)*)

- Optimal?
  - **Yes**
  - Also optimally efficient:
    - No other optimal algorithm will expand fewer nodes, for a given heuristic
- Time?
  - **Exponential** in worst case

- Space?
  - **Exponential** in worst case

- A* expands all nodes with f(n) < C*
  - This will be exponentially large

Recursive Best-First Search (RBFS)

- In practice A* runs out of memory before it runs out of time

  - How can we solve the memory problem for A* search?

  - **RBFS**

- Similar to DFS, but keeps track of the f-value of the best alternative path available from any ancestor of the current node

- If current node exceeds f-limit -> backtrack to alternative path

- As it backtracks, replace f-value of each node along the path with the best f(n) value of its children

  - This allows it to return to this subtree, if it turns out to look better than alternatives

Recursive Best First Search: Example

(a) After expanding Arad, Sibiu, Rimnicu Vilcea

- Path until **Rumnicu Vilcea** is already expanded

- **Above node**; *f*-limit for every recursive call is shown on top.

- Below node: *f(n)*

- The path is followed until Pitesti which has a *f*-value worse than the *f-limit*.

RBFS example

(b) After unwinding back to Sibiu and expanding Fagaras

- **Unwind recursion** and store best *f*-value for current best leaf Pitesti

  result, *f* [*best*] ← RBFS(*problem, best,* min(*f_limit, alternative*))

- *best* is now Fagaras. Call RBFS for new *best*
  - *best* value is now 450 which is larger than best f-value

RBFS example

(c) After switching back to Rimnicu Vilcea and expanding Pitesti
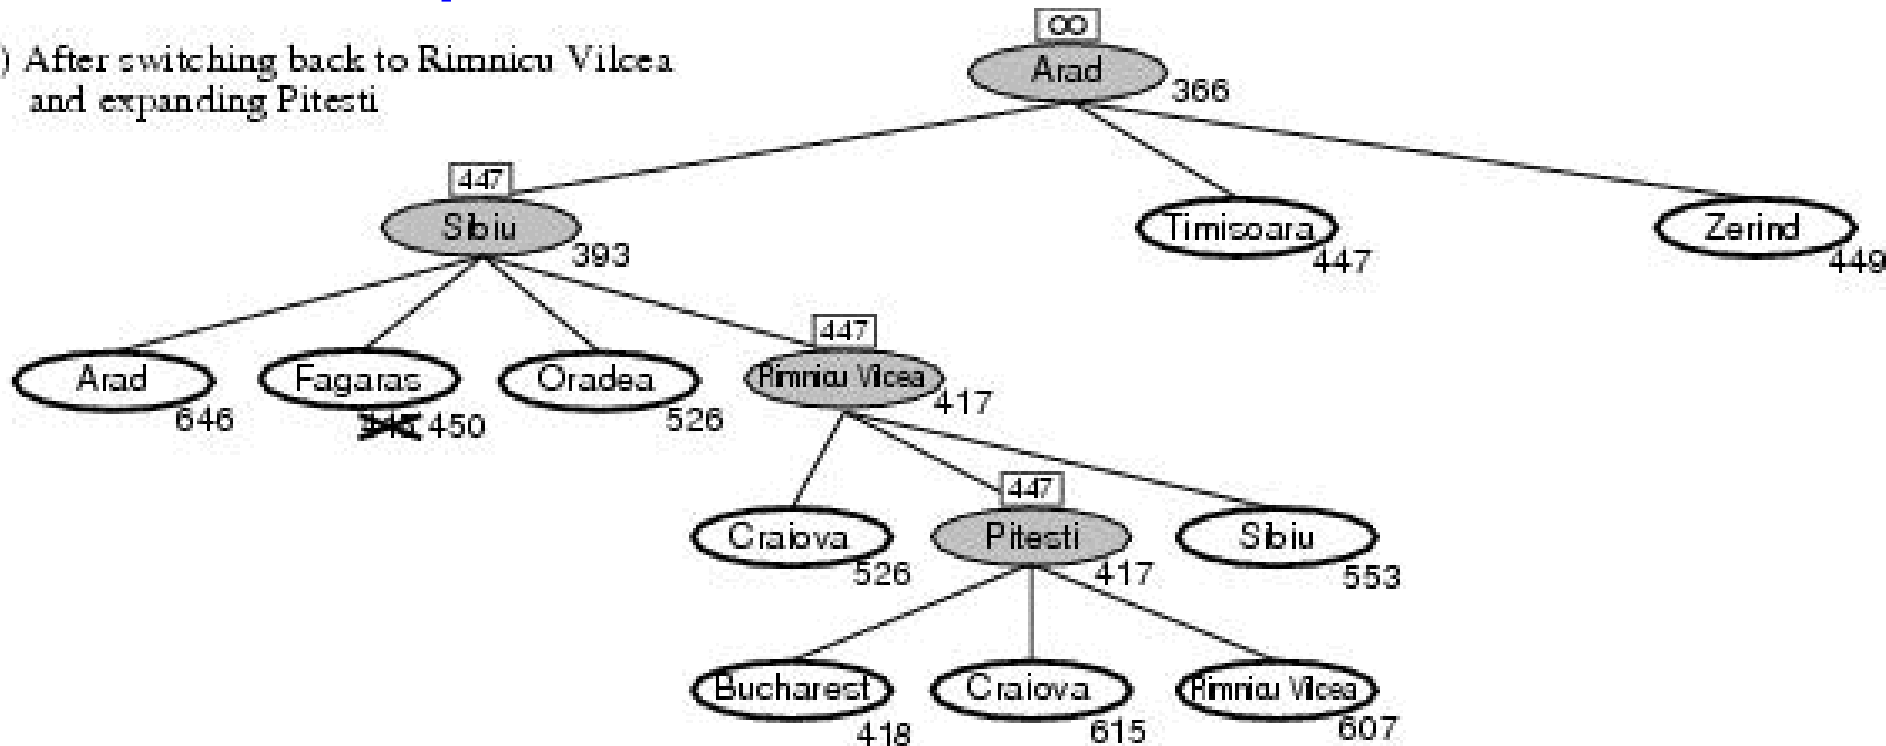
- **Unwind recursion again** and store best *f*-value for current best leaf Fagaras

    result, *f* [*best*] ← RBFS(*problem, best,* min(*f_limit, alternative*))

- *best* is now Rimnicu Viclea (again). Call RBFS for new *best*
  – Subtree is again expanded.
  – Best *alternative* subtree is now through Timisoara.
- Solution is found since because 447 > 418.

RBFS properties

- Like A*, optimal if *h(n)* is admissible

- Time complexity difficult to characterize
  – Depends on accuracy if h(n) and how often best path changes.
  – Can end up "switching" back and forth

- Space complexity is *O(bd)*
  – Other extreme to A* - uses ***too little*** memory.

(Simplified) Memory-bounded A* (SMA*)

- This is like A*, but when memory is full we delete the worst node (largest f-value).

- Like RBFS, we remember the best descendant in the branch we delete.

- If there is a tie (equal f-values) we delete the oldest nodes first.

- simplified-MA* finds the optimal *reachable* solution given the memory constraint.

- **Time can still be exponential**.

Designing heuristic functions

- Heuristics for the 8-puzzle

  $h_1(n)$ = number of misplaced tiles

  $h_2(n)$ = total Manhattan distance (number of squares from desired location of each tile)

Start State   Goal State

$$h_1(\text{start}) = 8$$
$$h_2(\text{start}) = 3+1+2+2+2+3+3+2 = 18$$

- Are h_1 and h_2 admissible?

Notion of dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible) then h_2 **dominates** h_1

  In other words h_2 is better for search

- Typical search costs (average number of nodes expanded) for 8-puzzle problem

  d=12   IDS = 3,644,035 nodes

  $A^*(h_1)$ = 227 nodes

  $A^*(h_2)$ = 73 nodes

  d=24   IDS = too many nodes

  $A^*(h_1)$ = 39,135 nodes

  $A^*(h_2)$ = 1,641 nodes

Effective branching factor

- **Effective branching factor b***

  – Is the branching factor that a uniform tree of depth *d* would have in order to contain *N+1* nodes.

  $$N+1 = 1 + b* + (b*)^2 + ... + (b*)^d$$

  – Measure is fairly constant for sufficiently hard problems.

    • Can thus provide a good guide to the heuristic's overall usefulness.

Effectiveness of different heuristics

	Search Cost			Effective Branching Factor		
d	IDS	A*(h_1)	A*(h_2)	IDS	A*(h_1)	A*(h_2)
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

- Results averaged over random instances of the 8-puzzle

Inventing heuristics via "relaxed problems"

- A problem with **fewer restrictions** on the actions is called a relaxed problem

- **<u>The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem</u>**

- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then $h_1(n)$ gives the shortest solution

- If the rules are relaxed so that a tile can move to any adjacent square, then $h_2(n)$ gives the shortest solution

Summary

- Uninformed search methods have their limits

- Informed (or heuristic) search uses problem-specific heuristics to improve efficiency
  - Best-first
  - A*
  - RBFS
  - SMA*
  - Techniques for generating heuristics

- Can provide significant speed-ups in practice
  - e.g., on 8-puzzle
  - But can still have worst-case exponential time complexity